

Jordan Rivera

Brooklyn, NY | (555) 010-4477 | jordan.rivera@example.com | [GitHub](#) | [LinkedIn](#)

September 23, 2026

Dear Hiring Team,

Proseware is building the part of AI I most want to work on: deciding when people can trust what an agent did. The Senior Full-Stack Engineer role on Agent Workspace, and especially "the approval flows that decide what an agent may do on its own," describes the problem I have spent the last two years solving at a smaller scale.

At Contoso Payments I lead a 6-engineer squad that owns the Merchant Dashboard, used by more than 40,000 small-business merchants. I also built and own our LLM-assisted support triage service on the Claude API. It uses tool calling to pull a merchant's account, payout, and dispute status, then drafts a reply for a support agent to review. Sixty agents use it daily, and median first-response time fell 31%. The part I am proudest of is the unglamorous part: a 200-ticket evaluation set that runs on every prompt or model change, and a human review step on every reply. Before that, I led our dispute flow redesign, starting with three rounds of merchant interviews, and cut average merchant response time from 6 days to 2.

I build the same way outside work. Pantry Planner, my meal-planning web app with about 1,200 monthly active users, generates plans on the Claude API with structured JSON output and checks every plan against allergy and diet rules before anyone sees it. Earlier, at Fabrikam Health, I added audit logging for patient-record access that closed the findings from a HIPAA security review, so I know what building for people who carry real risk demands.

I would welcome a conversation about making Proseware's agents easier to trust. I am based in Brooklyn and glad to be in the New York office.

Best,

Jordan Rivera